



Deep Learning

Deep Feedforward Networks

Lecturer: Duc Dung Nguyen, PhD.

Contact: nddung@hcmut.edu.vn

Faculty of Computer Science and Engineering
Hochiminh city University of Technology

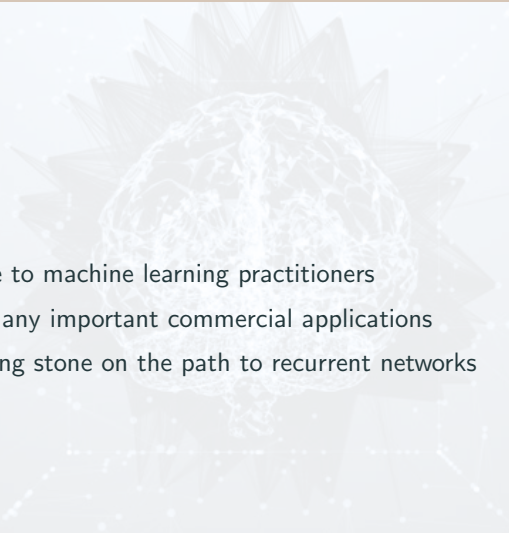
- 
- A large, faint background graphic of a human brain with a network of white lines representing neural connections. The brain is centered and surrounded by a starburst-like pattern of lines.
1. Deep Networks
 2. Gradient Based Learning
 3. Hidden Units
 4. Architecture Design
 5. Back-Propagation

Deep Networks



Deep feedforward networks (multilayer perceptrons (MLPs))

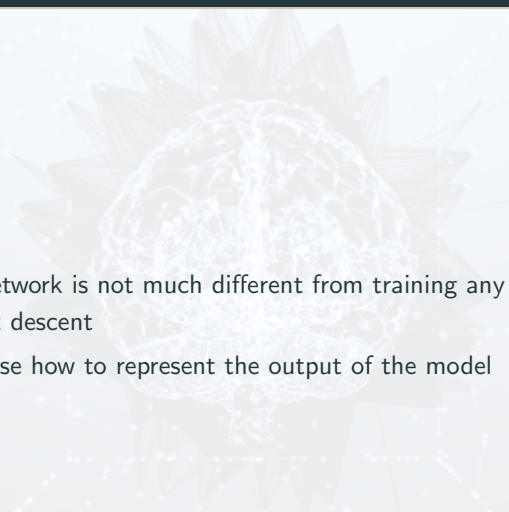
- The quintessential deep learning models
- Goal: approximate some function f^*
- Information flow through the function being evaluated
- No feedback connection

- 
- A faint, stylized background graphic of a lotus flower is centered on the slide. The lotus is composed of many small, interconnected nodes and lines, giving it a network-like appearance. The nodes are represented by small white dots, and the lines are thin white lines connecting the dots. The lotus has multiple layers of petals, each layer made of these network-like structures. The overall color of the background graphic is a light gray/white, blending into the light gray background of the slide.
- Extreme importance to machine learning practitioners
 - Form the basis of many important commercial applications
 - A conceptual stepping stone on the path to recurrent networks

Linear models

- Logistic regression, linear regression
- Can be fit efficiently and reliably
- Can obtain closed form solution or with convex optimization
- Limitation: capacity is limited to linear functions
 - Can not understand the interaction between any two input variables

Gradient Based Learning

- 
- A large, faint, stylized graphic of a lotus flower is centered in the background. The lotus is composed of many small, interconnected nodes and lines, giving it a digital or network-like appearance. The nodes are small circles, and the lines are thin, creating a complex web-like structure that forms the petals and center of the flower.
- Training a neural network is not much different from training any other machine learning model with gradient descent
 - Cost function: choose how to represent the output of the model

Cost function

- More or less the same as those for other parametric models, such as linear models
- The total cost function used to train a neural network will often combine **one of the primary cost functions** with a **regularization term**

- Most modern neural networks are trained using maximum likelihood
- The cost function is simply the negative log-likelihood, equivalently described as the cross-entropy between the training data and the model distribution.

$$J(\theta) = -\mathbb{E}_{x,y \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{y}|\mathbf{x}). \quad (1)$$

- Cost function changes from model to model, depending on the specific form of $\log p_{\text{model}}$
- If $p_{\text{model}}(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}; f(\mathbf{x}; \theta), \mathbf{I})$ then

$$J(\theta) = \frac{1}{2} \mathbb{E}_{x,y \sim \hat{p}_{\text{data}}} \|\mathbf{y} - f(\mathbf{x}; \theta)\|^2 + C$$

- A sufficiently powerful neural network: be able to represent any function from a wide class of functions
- **Learning**: choosing a function rather than merely choosing a set of parameters
- **Mean squared error** and **mean absolute error** often lead to poor results when used with *gradient-based optimization*
 - Some output units that saturate produce very small gradients when combined with these cost functions.
 - **Cross-entropy** cost function is more popular

1. Linear Units for Gaussian Output Distributions

- Given features \mathbf{h} , a layer of linear output units produces a vector $\hat{\mathbf{y}} = \mathbf{W}^\top \mathbf{h} + b$
- The mean of a conditional Gaussian distribution:

$$p(y|\mathbf{x}) = \mathcal{N}(y; \hat{\mathbf{y}}, \mathbf{I}). \quad (2)$$

- Maximizing the log-likelihood is then equivalent to minimizing the mean squared error

2. Sigmoid Units for Bernoulli Output Distributions

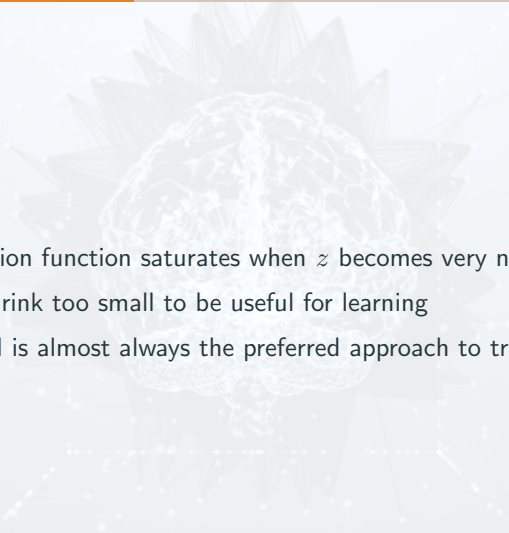
- Task: predicting the value of a binary variable y
- The neural net needs to predict only $P(y = 1|x)$
- **What if:**

$$P(y = 1|x) = \max \{0, \min \{1, \mathbf{w}^\top \mathbf{h} + b\}\} . \quad (3)$$

- It is better to ensure that there is always a strong gradient whenever the model has the wrong answer
- Sigmoid output:

$$\hat{y} = \sigma(\mathbf{w}^\top \mathbf{h} + b) \quad (4)$$

- We may see this output as a combination of linear transformation $z = \mathbf{w}^\top \mathbf{h} + b$ and an activation function $\sigma(z)$

- 
- A faint, stylized background graphic of a lotus flower is centered on the slide. The lotus is composed of many small, interconnected nodes and lines, giving it a digital or network-like appearance. The nodes are small circles, and the lines are thin, creating a complex web-like structure that forms the petals and stem of the flower.
- The sigmoid activation function saturates when z becomes very negative or very positive
 - The gradient can shrink too small to be useful for learning
 - Maximum likelihood is almost always the preferred approach to training sigmoid output units

3. Softmax Units for Multinoulli Output Distributions

- **Softmax functions:** are most often used as the output of a classifier, to represent the probability distribution over n different classes
- Linear layer predicts unnormalized log probability:

$$\mathbf{z} = \mathbf{W}^\top \mathbf{h} + \mathbf{b} \quad (5)$$

- Softmax function:

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)} \quad (6)$$

- Many objective functions other than the log-likelihood do not work as well with the softmax function (e.g. squared error)
 - Vanishing gradient

- Log-likelihood can undo the exp of the softmax

$$\log \text{softmax}(\mathbf{z})_i = z_i - \log \sum_j \exp(z_j) \quad (7)$$

- Softmax output is invariant to adding scalar

$$\begin{aligned} \text{softmax}(\mathbf{z}) &= \text{softmax}(\mathbf{z} + c) \\ \text{softmax}(\mathbf{z}) &= \text{softmax}(\mathbf{z} - \max_i z_i) \end{aligned} \quad (8)$$

Hidden Units



- How to choose the type of hidden units

- **Rectified linear units (ReLU)**: use activation function $g(z) = \max\{0, z\}$
- The gradient is useful for learning (no second-order effect)
- ReLU is typically used on top of an affine transformation

$$\tilde{h} = g(\mathbf{W}^\top \mathbf{x} + \mathbf{b}) \quad (9)$$

- Initialization is important!

- **Drawback:** cannot learned via gradient-based methods on examples for which their activation is 0
- Generalization:

$$h_i = g(\mathbf{z}, \alpha)_i = \max(0, z_i) + \alpha_i \min(0, z_i) \quad (10)$$

- Absolute value rectification: fix $\alpha_i = -1$
- Leaky ReLU (Maas et al., 2013)
- Parametric ReLU (PReLU) (He et al., 2015)

- **Maxout units (Goodfellow et al., 2013):** generalize ReLU
 - Divide \mathbf{z} into groups of k values
 - Each maxout unit outputs the maximum element of one of these groups

$$g(\mathbf{z})_i = \max_{j \in \mathbb{G}^{(i)}} z_j \quad (11)$$

- $\mathbb{G}^{(i)}$ is the set of indices for group i
- A maxout unit can learn a piecewise linear, convex function with up to k pieces
- *Learning the activation function itself*

- Hyperbolic tangent activation function:

$$g(z) = \tanh(z) = 2\sigma(2z) - 1 \quad (12)$$

- The widespread saturation of sigmoid unit can make gradient-based learning very difficult
- Tangent activation function typically performs better than logistic sigmoid (resemble the identity function more closely)

Architecture Design

- The **architecture** refers to the overall structure of the network:
 - How many units it should have
 - How these units should be connected to each other
- Most NN are organized into groups of units called **layers**
- *Chain structure*

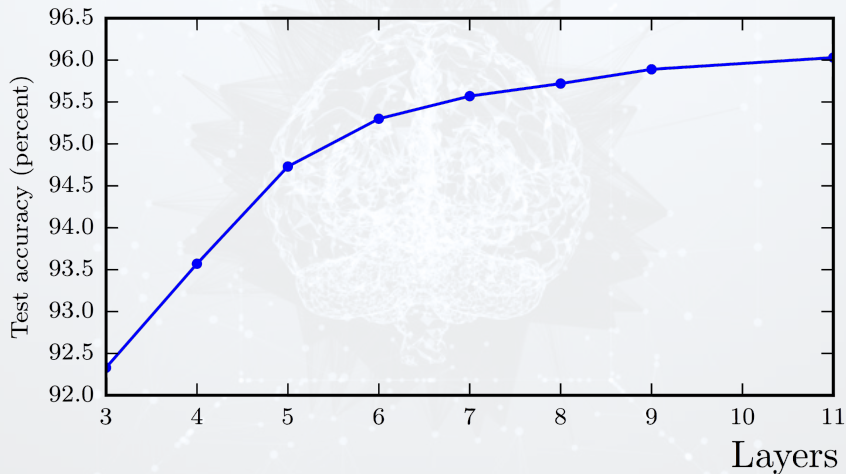
- **Linear model:** represent only linear functions.
 - Easy to train: many loss functions result in convex optimization problems when applied to linear models
- **The universal approximation theorem:** regardless of what function we are trying to learn, we know that a large MLP will be able to represent this function.
 - We are not guaranteed that the training algorithm will be able to learn that function


Learning can fail for two different reasons

- The optimization algorithm used for training may not be able to find the value of the parameters that corresponds to the desired function
- The training algorithm might choose the wrong function due to overfitting

Depth

- A **feedforward network** with a **single layer** is sufficient to represent any function
 - The layer may be **infeasibly large** and may fail to learn and generalize correctly
- Deeper models can *reduce the number of units* required to represent the desired function and can *reduce the amount of generalization error*







Back-Propagation





- **Feed-forward neural network:** information flows forward through the network.
- **Forward propagation:** the inputs x provide the initial information that then propagates up to the hidden units at each layer and finally produces \hat{y} .
- **The back-propagation algorithm:** allows the information from the cost flow backwards through the network, in order to compute the gradient.



(a)



(b)



Chain Rule of Calculus Let x be a real number, and let f and g both be functions mapping from a real number to a real number. Suppose that $y = g(x)$, $z = f(g(x)) = f(y)$.

- The chain rule

$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx} \quad (13)$$

- Generalization: $\mathbf{x} \in \mathbb{R}^m$, $\mathbf{y} \in \mathbb{R}^n$, $g : \mathbb{R}^m \rightarrow \mathbb{R}^n$, and $f : \mathbb{R}^n \rightarrow \mathbb{R}$

$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i} \quad (14)$$

- **Vector notation:**

$$\Delta_x \mathbf{z} = \left(\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \right)^\top \Delta_y \mathbf{z} \quad (15)$$

where $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ is the $n \times m$ Jacobian matrix of g

- The gradient of a variable \mathbf{x} can be obtained by multiplying a Jacobian matrix $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$ by a gradient $\Delta_y \mathbf{z}$

