



# Deep Learning

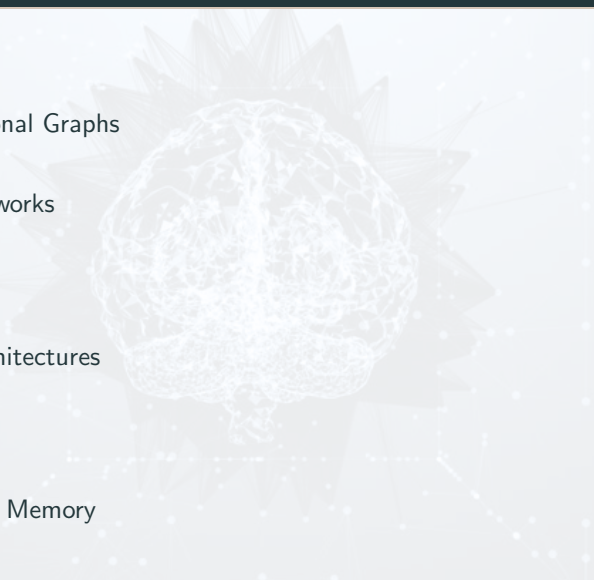
## Recurrent Neural Networks

---

Lecturer: Duc Dung Nguyen, PhD.

Contact: [nddung@hcmut.edu.vn](mailto:nddung@hcmut.edu.vn)

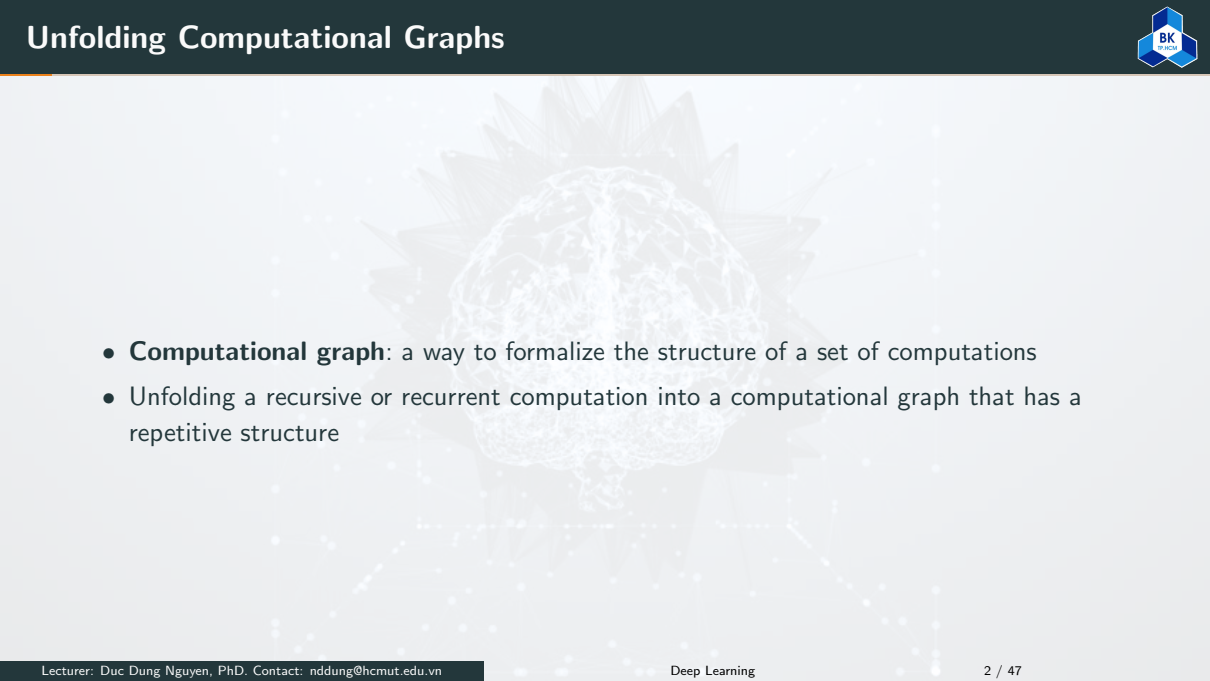
Faculty of Computer Science and Engineering  
Hochiminh city University of Technology

- 
- A large, faint, stylized graphic of a brain with neural connections is centered in the background. It is overlaid with a network of white dots and lines, suggesting a computational graph or neural network structure.
1. Unfolding Computational Graphs
  2. Recurrent Neural Networks
  3. Bidirectional RNNs
  4. Encoder-Decoder Architectures
  5. Other RNNs
  6. The Long Short-Term Memory



# Unfolding Computational Graphs


---

- 
- The background of the slide features a large, faint image of a lotus flower. Overlaid on the lotus is a complex, glowing white network of nodes and lines, resembling a neural network or a computational graph, which adds a technical and artistic touch to the presentation.
- **Computational graph:** a way to formalize the structure of a set of computations
  - Unfolding a recursive or recurrent computation into a computational graph that has a repetitive structure

- The classical form of a dynamical system

$$s^{(t)} = f(s^{(t-1)}; \theta) \quad (1)$$


where  $s^{(t)}$  is called the state of the system




- Any function involving recurrent can be considered as a feedforward network
- A single, shared model allows generalization to sequence lengths that did not appear in the training set
- Allows the model to be estimated with far fewer training examples
- The unfolded graph provides an explicit description of which computations to perform

# Recurrent Neural Networks


---








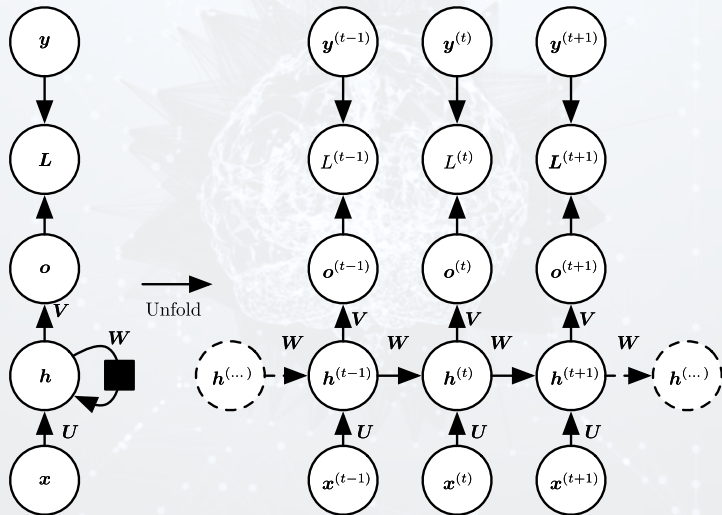
Vanilla Neural Networks

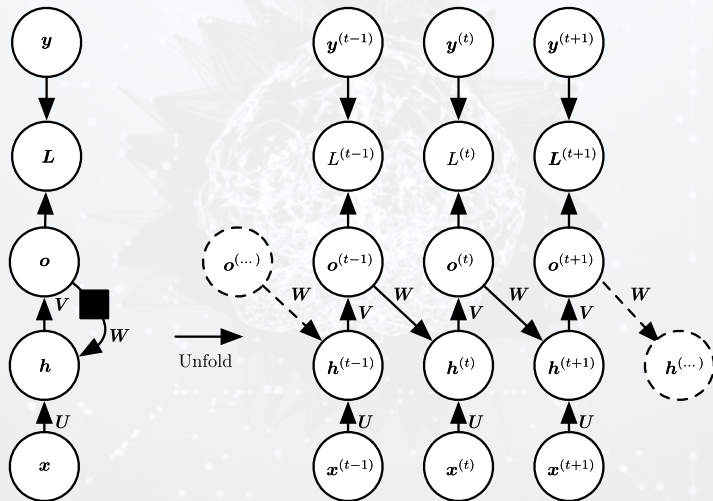





Some examples of important design patterns:

- RNNs that produce an output at each time step and have recurrent connections between hidden units
- RNNs that produce an output at each time step and have recurrent connections only from the output at one time step to the hidden units at the next time step
- RNNs with recurrent connections between hidden units, that read an entire sequence and then produce a single output







- Less powerful: lacks hidden-to-hidden recurrent connections
- Cannot simulate a universal Turing machine
- The output units are explicitly trained to matched the training set targets → unlikely to capture the necessary information about the past history
- Describe full state of the system?



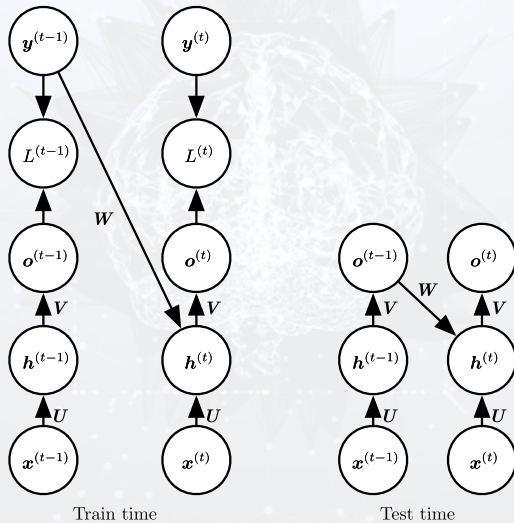
## Teacher forcing

- A procedure that emerges from the **maximum likelihood criterion**: during training, the model receives the ground truth output  $y(t)$  as input at time  $t + 1$ .
- The conditional maximum likelihood criterion is

$$\log p \left( y^{(1)}, y^{(2)} | x^{(1)}, x^{(2)} \right) = \log p \left( y^{(2)} | y^{(1)}, x^{(1)}, x^{(2)} \right) + \log p \left( y^{(1)} | x^{(1)}, x^{(2)} \right) \quad (2)$$

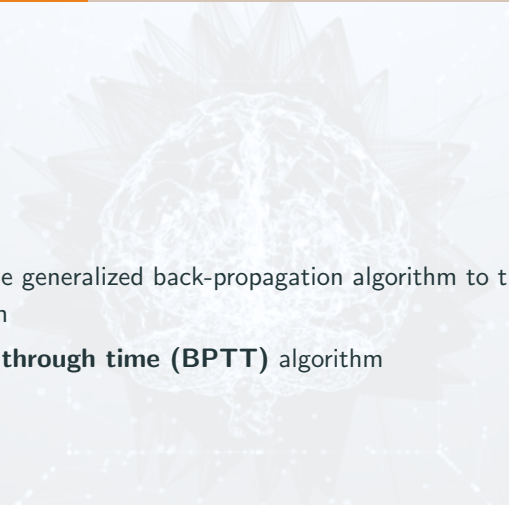
- During training: feeding the model's own output back into itself  $\rightarrow$  these connections should be fed with the target values specifying what the correct output should be

# Teacher Forcing and Networks with Output Recurrence



## Disadvantages

- When the network is going to be used in an open-loop mode → problem!
- Inputs in training can be quite different from the test time

- 
- A faint, stylized background graphic of a brain with neural connections, rendered in a light gray color, is centered behind the text.
- **Training:** applies the generalized back-propagation algorithm to the unrolled computational graph
  - **Back-propagation through time (BPTT)** algorithm

- The nodes of our computational graph include the parameters  $U, V, W, b$  and  $c$
- The sequence of nodes are indexed by  $t$  for  $x^{(t)}, h^{(t)}, o^{(t)}$  and  $L^{(t)}$
- For each node  $N$  we need to compute the gradient  $\nabla_N L$  recursively
- Start the recursion with the nodes immediately preceding the final loss

$$\frac{\partial L}{\partial L^{(t)}} = 1 \quad (3)$$

- Assume:
  - Outputs  $o^{(t)}$  are used as the argument to the softmax function to obtain the vector  $\hat{y}$  of probabilities over the output
  - The loss is the negative log-likelihood of the true target  $y^{(t)}$
- The gradient  $\nabla_{o^{(t)}} L$  on the outputs at time step  $t$ , for all  $i, t$ :

$$(\nabla_{o^{(t)}} L)_i = \frac{\partial L}{\partial o_i^{(t)}} = \frac{\partial L}{\partial L^{(t)}} \frac{\partial L^{(t)}}{\partial o_i^{(t)}} = \hat{y}_i^{(t)} - 1_{i, y^{(t)}} \quad (4)$$

## Back-propagation

- Starting from the end of the sequence
- At the final time step  $\tau$ ,  $h^{(\tau)}$  only has  $o^{(\tau)}$  as a descendent, so its gradient is simple:

$$\begin{aligned}\nabla_{h^{(t)}} L &= \left( \frac{\partial h^{(t+1)}}{\partial h^{(t)}} \right)^\top (\nabla_{h^{(t+1)}} L) + \left( \frac{\partial o^{(t)}}{\partial h^{(t)}} \right)^\top (\nabla_{o^{(t)}} L) \\ &= W^\top (\nabla_{h^{(t+1)}} L) \text{diag} \left( 1 - \left( h^{(t+1)} \right)^2 \right) + V^\top (\nabla_{o^{(t)}} L)\end{aligned}\tag{5}$$


- How about  $\nabla_W f$ ?
- Introduce dummy variables  $W^{(t)}$  that are defined to be copies of  $W$  but with each  $W^{(t)}$  used only at time step  $t$
- Use  $\nabla_{W^{(t)}}$  to denote the contribution of the weights at time step  $t$  to the gradient

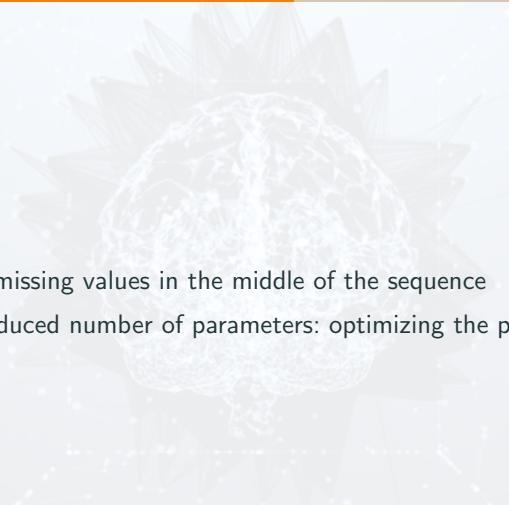


The gradient on the remaining parameters

$$\begin{aligned}\nabla_c L &= \sum_t \left( \frac{\partial o^{(t)}}{\partial c} \right)^\top \nabla_{o^{(t)}} L = \sum_t \nabla_{o^{(t)}} L \\ \nabla_b L &= \sum_t \left( \frac{\partial h^{(t)}}{\partial b^{(t)}} \right)^\top \nabla_{h^{(t)}} L = \sum_t \text{diag} \left( 1 - \left( h^{(t)} \right)^2 \right) \nabla_{h^{(t)}} L \\ \nabla_V L &= \sum_t \sum_i \left( \frac{\partial L}{\partial o_i^{(t)}} \right) \nabla_V o_i^{(t)} = \sum_t (\nabla_{o^{(t)}} L) h^{(t)\top}\end{aligned}\tag{6}$$

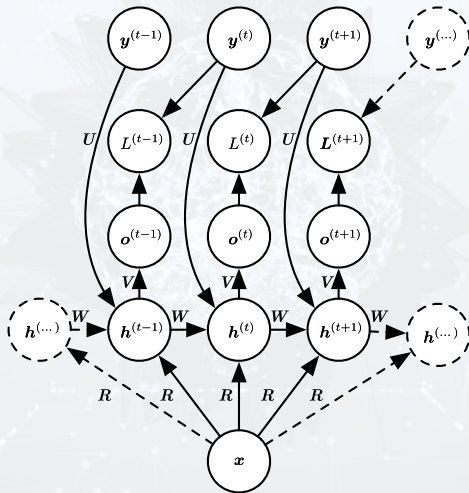
$$\begin{aligned}\nabla_W L &= \sum_t \sum_i \left( \frac{\partial L}{\partial h_i^{(t)}} \right) \nabla_{W(t)} h^{(t)} \\ &= \sum_t \text{diag} \left( 1 - \left( h^{(t)} \right)^2 \right) (\nabla_{h^{(t)}} L) h^{(t-1)\top} \\ \nabla_U L &= \sum_t \sum_i \left( \frac{\partial L}{\partial h_i^{(t)}} \right) \nabla_{U(t)} h^{(t)} \\ &= \sum_t \text{diag} \left( 1 - \left( h^{(t)} \right)^2 \right) (\nabla_{h^{(t)}} L) x^{(t)\top}\end{aligned}\tag{7}$$

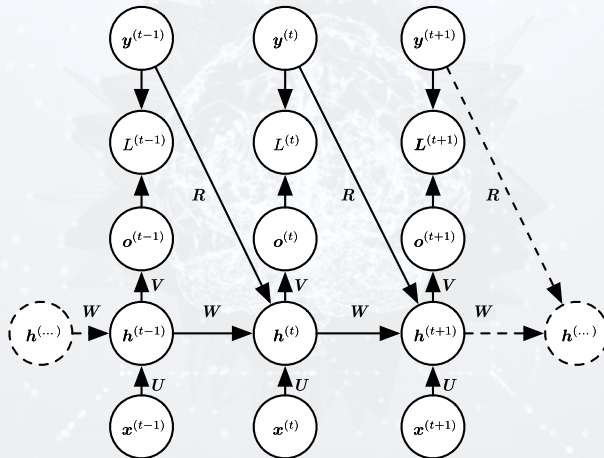


- 
- A faint, stylized background image of a human brain with a network of white lines representing neural connections or data flow, set against a light gray background.
- Difficult to predict missing values in the middle of the sequence
  - The price for the reduced number of parameters: optimizing the parameters may be difficult

## Some common ways of providing an extra input to an RNN:

- As an extra input at each time step
- As the initial state  $h^{(0)}$
- Both

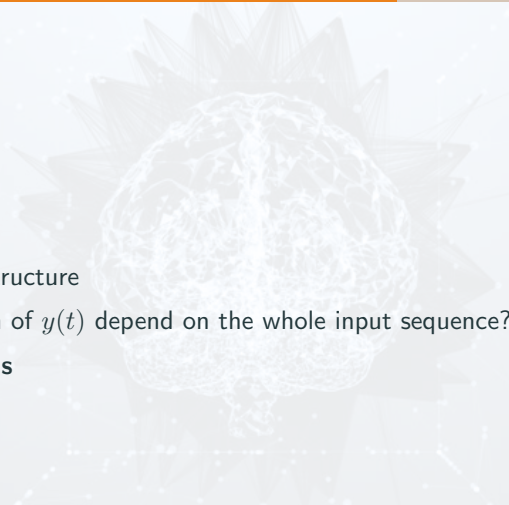


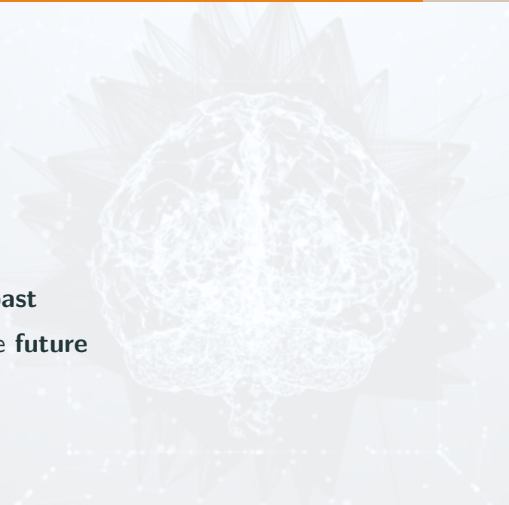



# Bidirectional RNNs

---



- 
- A faint, stylized background graphic of a human brain with a network of white lines representing neural connections or data flow, set against a light gray background.
- So far: “**causal**” structure
  - What if a prediction of  $y(t)$  depend on the whole input sequence?
  - **Bidirectional RNNs**

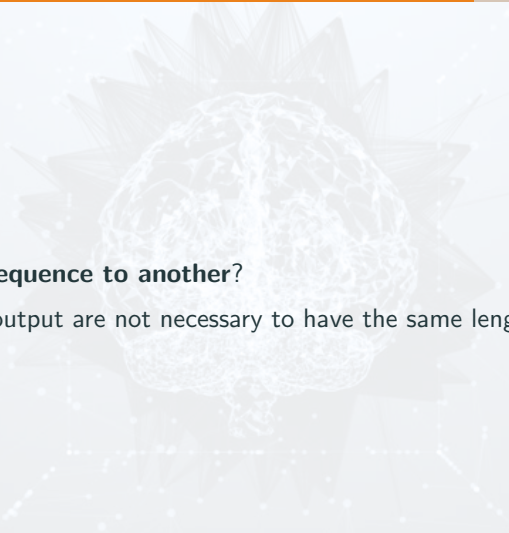
- 
- A faint, stylized background graphic of a human brain with a network of white lines representing neural connections. The brain is centered and has a spiky, sun-like border. The background is a light gray with a subtle grid of dots and lines.
- Let remember the **past**
  - And think about the **future**





# Encoder-Decoder Architectures

---

- 
- A faint, stylized background graphic of a human brain with a network of white lines representing neural connections or data flow, set against a light gray background.
- Can we **map one sequence to another**?
  - The input and the output are not necessary to have the same length
  - **Examples?**


## Examples

- Speech recognition
- Machine translation
- Question answering
- Etc.



- Input to RNN: **context**
- Produce a representation of context  $C$
- **Context**  $C$ : a vector or sequence of vectors that summarize the input sequence  $X = (x^{(1)}, \dots, x^{(n_x)})$

# Encoder-Decoder Sequence-to-Sequence Architectures






- Encoder and decoder do not need to have the same size
- **Limitation:** the output of encoder has a dimension that is too small to properly summarize a long sequence


## Other RNNs

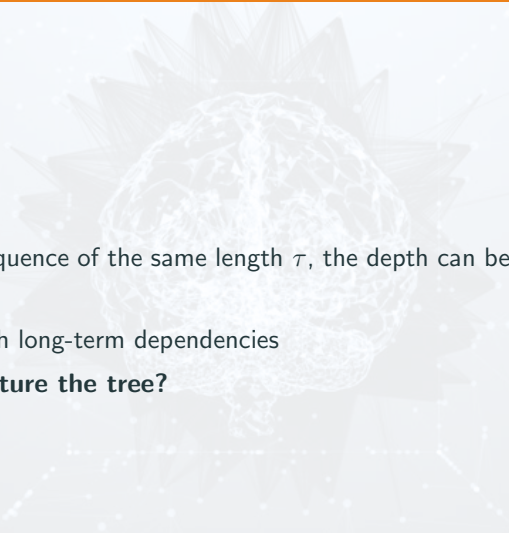
---



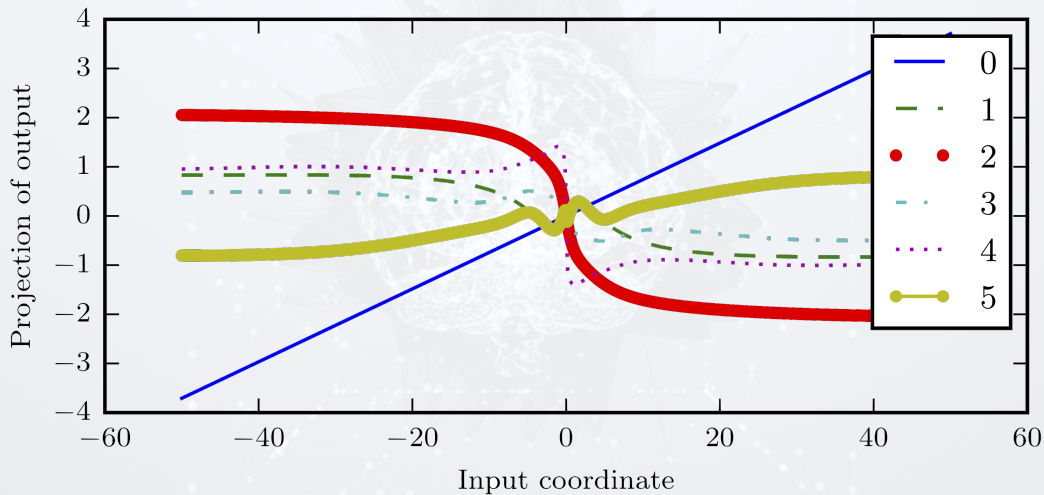


- A generalization of recurrent networks
- A different kind of computational graph: a deep tree, rather than the chain-like structure of RNNs
- E.g.: processing data structures as input to neural nets (both in NLP as well as in CV)



- 
- A faint, stylized diagram of a Recursive Neural Network (RNN) tree structure is visible in the background. It shows a central node branching out into multiple children, which then further branch out, illustrating the hierarchical nature of the network. The diagram is composed of light gray lines and dots, giving it a network-like appearance.
- Advantage: for a sequence of the same length  $\tau$ , the depth can be drastically reduced from  $\tau$  to  $O(\log \tau)$
  - Might help deal with long-term dependencies
  - **How to best structure the tree?**

# The Challenge of Long-term Dependencies





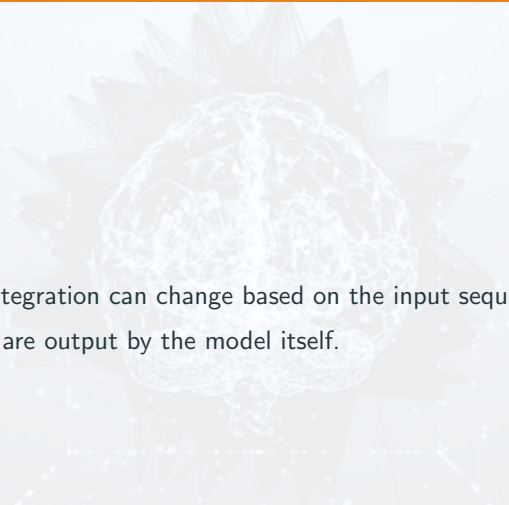
# **The Long Short-Term Memory**

---




## LSTM (Hochreiter and Schmidhuber, 1997)

- **LSTM**: introduce self-loops to produce paths where gradient can flow for a long durations.
- Make the **weight** on this self-loop **conditioned on the context**, rather than fixed
- **Gates**
- Time scale of integration can be changed dynamically

- 
- A faint, stylized background graphic of a brain with neural connections, overlaid on a grid of dots and lines.
- The time scale of integration can change based on the input sequence
  - The time constants are output by the model itself.

# The Long Short-Term Memory



- Forget gate

$$f_i^{(t)} = \sigma \left( b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} \right) \quad (8)$$

- Internal state

$$s_i^{(t)} = f_i^{(t)} s_i^{(t-1)} + g_i^{(t)} \sigma \left( b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right) \quad (9)$$

- External input gate

$$g_i^{(t)} = \sigma \left( b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} \right) \quad (10)$$

- Output gate  $q_i^{(t)}$

$$h_i^{(t)} = \tanh \left( s_i^{(t)} \right) q_i^{(t)} \quad (11)$$

$$q_i^{(t)} = \sigma \left( b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} \right) \quad (12)$$

- LSTM networks learns long-term dependencies better
- Optimization
  - Clipping gradient
  - Regularizing: encourage information flow
- Case studies:
  - Memory networks (Westion et al., 2014)
  - Neural Turing machine (Graves et al., 2014)
  - Multiple object recognition with attention (Ba et al.)
  - Image captioning